

What it does `prefetch_summary_queue` builds a limited list of pre-fetch jobs from a set of file changes. It first filters the changes by `.gitignore` rules, reserved "novadiff" documents, and non-text files. After filtering, it computes a subtree weight for each path, sorts the eligible changes by depth, weight, and path, and then takes up to `limit` (but at least one) to create `PrefetchSummaryJob` entries. The function returns a `PrefetchSummaryQueueOut` that includes the queue, counts of input and skipped changes, and how many jobs were actually returned.

How it works
1. Initial counts – `input_changes` records the total number of changes.
2. **Gitignore filtering*

- if `gitignore_roots` is provided, `filter_changes_for_prefetch_gitignore` is called; otherwise all changes are kept.
- 3. **Reserved docs filtering*
 - `filter_out_novadiff_docs_changes` removes any "novadiff" documents.
- 4. **Non-text filtering*
 - again conditioned on `gitignore_roots`, `filter_out_non_text_changes` removes binary files.
- 5. **Weight calculation*
 - `subtree_change_counts` tallies how many changes exist under each parent directory.
- 6. **Sorting*
 - eligible changes are sorted by path depth, then by the weight of their parent subtree, and finally lexicographically.
- 7. **Queue construction*
 - the first `cap` (max of `limit` and 1) sorted changes are mapped to `PrefetchSummaryJob` structs, capturing the original path, kind, depth, and parent subtree weight.
- 8. **Return*
 - a `PrefetchSummaryQueueOut` is assembled with all counts and the job queue.

Notes for reviewers- The use of `unwrap_or(&0)` assumes every parent directory will have a weight entry; if `subtree_change_counts` can omit a key, this may silently default to zero.- `limit.max(1)` guarantees at least one job, even if no changes remain after filtering; confirm this is intended behavior.- Path normalization replaces backslashes with slashes; ensure this matches the rest of the codebase's path handling.- The sorting logic prioritizes depth before weight; consider whether this ordering aligns with pre-fetch performance goals.- The function returns counts of skipped items; verify that these metrics are used elsewhere for diagnostics.