

packages/graph-view/scripts/benchmark-aggregations.mjs

Kind: added · Baseline: a9c34bfb03fe · Target: NovaDiff · Generated: 2026-05-20T05:07:26.220Z

Overview

A new benchmark script `packages/graph-view/scripts/benchmark-aggregations.mjs` (diff R1-99) has been added to measure the performance difference between the legacy and refactored graph-aggregation logic used in ``use`

Overview

Graph`.

Key changes

- Import added: `import { performance } from "node:perf_hooks";` (R16).
- Graph construction helper `makeGraph(layerCount, nodesPerLayer)` (lines 18-31).
- Legacy aggregation `aggregateBefore(graph)` (lines 35-49).
- Optimized aggregation `aggregateAfter(graph, nodesById)` (lines 53-72).
- Benchmark runner `bench(label, layerCount, nodesPerLayer)` (lines 75-93) times both algorithms, logs speed-up and parity.
- Four benchmark runs executed at the bottom of the file (lines 96-99).

Impact

The script outputs console logs with timing, speed-up factor, and parity check, providing a way to verify correctness and performance of the aggregation logic without altering production code.

Risks & follow-ups

- Node compatibility: relies on `node:perf_hooks`; ensure the target runtime supports it.
- Benchmark reliability: timings can vary with CPU load and V8 optimizations; consider multiple runs for precise numbers.
- Future refactors: if aggregation logic changes again, update the benchmark accordingly.
- Documentation: add instructions on running the script and interpreting results for future maintainers.